

FindDefault-Credit card fraud detection

Introduction

Problem Statement:

A credit card is one of the most used financial products to make online purchases and payments. Though the Credit cards can be a convenient way to manage your finances, they can also be risky. Credit card fraud is the unauthorized use of someone else's credit card or credit card information to make purchases or withdraw cash.

It is important that credit card companies are able to recognize fraudulent credit card transactions so that customers are not charged for items that they did not purchase.

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

We have to build a classification model to predict whether a transaction is fraudulent or not.

Data Exploration

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...	V21	V22	V23	V24	V25	V26	V27	V28	Class	Scaled_Amount
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098898	0.363787	0.090794	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	0	0.244964
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	-0.166974	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	0	-0.342475
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	0.207643	...	0.247998	0.771679	0.909412	-0.689281	-0.327842	-0.139097	-0.055353	-0.059752	0	1.160686
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	-0.054952	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	0	0.140534
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	0.753074	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	0	-0.073403
...
284802	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	4.356170	...	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0	-0.350151
284803	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	-0.975926	...	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	0	-0.254117
284804	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	-0.484782	...	0.232045	0.578229	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561	0	-0.081839
284805	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	-0.399126	...	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	0	-0.313249
284806	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	-0.915427	...	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	0	0.514355

284807 rows x 30 columns

To protect the user's identity and the security of their confidential information, the dataset provider has applied Principal Component Analysis transformation on the original numerical features and compressed it into 28 principals components.

Only two features have not been transformed i.e. 1) Time and 2) Amount

The feature class will be target column with user labels as:

0 : non-fraudulent 1 : fraudulent

FindDefault-Credit card fraud detection


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284887 entries, 0 to 284886
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   Time        284887 non-null float64
1   V1          284887 non-null float64
2   V2          284887 non-null float64
3   V3          284887 non-null float64
4   V4          284887 non-null float64
5   V5          284887 non-null float64
6   V6          284887 non-null float64
7   V7          284887 non-null float64
8   V8          284887 non-null float64
9   V9          284887 non-null float64
10  V10         284887 non-null float64
11  V11         284887 non-null float64
12  V12         284887 non-null float64
13  V13         284887 non-null float64
14  V14         284887 non-null float64
15  V15         284887 non-null float64
16  V16         284887 non-null float64
17  V17         284887 non-null float64
18  V18         284887 non-null float64
19  V19         284887 non-null float64
20  V20         284887 non-null float64
21  V21         284887 non-null float64
22  V22         284887 non-null float64
23  V23         284887 non-null float64
24  V24         284887 non-null float64
25  V25         284887 non-null float64
26  V26         284887 non-null float64
27  V27         284887 non-null float64
28  V28         284887 non-null float64
29  Amount      284887 non-null float64
30  Class       284887 non-null int64  
dtypes: float64(30), int64(1)
```

We will get the information about the columns in the DataFrame we have created as :

The dataset exclusively comprises numerical features, and notably, there are no instances of missing values. Consequently, there is no need for null-value handling in this dataset.

Exploratory Data Analysis

- For the subsequent step, we will conduct fundamental Exploratory Data Analysis (EDA) on the dataset to enhance our understanding and extract valuable insights.



- The barplot reveals a significant imbalance between classes (0-Non Fradulent) and (1-Fraudulent).
- Majority of features are in PCA form , with the exceptions being Time and Amount ,a more in-depth examination of these two features is required.

FindDefault-Credit card fraud detection

```
df['Amount'].describe()
```

```
count    284807.000000
mean       88.349619
std       250.120109
min         0.000000
25%        5.600000
50%       22.000000
75%       77.165000
max      25691.160000
Name: Amount, dtype: float64
```

```
df['Time'].describe()
```

```
count    284807.000000
mean     94813.859575
std     47488.145955
min         0.000000
25%     54201.500000
50%     84692.000000
75%    139320.500000
max    172792.000000
Name: Time, dtype: float64
```

- We will initiate the visualization of transaction counts across hours, starting with the entire dataset. Subsequently, we will partition the dataset into fraudulent and non-fraudulent transactions to gain a more detailed perspective.
- Now we will check the number of occurrences of each class label and we will plot the information using matplotlib.

```
Number of Non-Fraudulent Transactions: 284315
Number of Fraudulent Transactions: 492
Percentage of Fraudulent Transactions: 0.17
```

- We can observe that the genuine transactions are over 99%.
 - We will apply scaling techniques on the "Amount" feature to transform the range of values.
 - We will drop the original "Amount" column and add a new column with the scaled values. We will also drop the "Time" columns as it is irrelevant.
-
- Now, we will split the credit card data with a split of 70-30 using train_test_split().
 - train_test_split() function in scikit-learn is a useful utility for splitting a dataset into training and testing sets.
 - Parameters
 - X: Feature matrix
 - Y: Target variable
 - test_size: Proportion of the dataset to include in the test split. Here we have set the test_size as 0.3 means 30% of the data we take as testing data set.
 - random_state: we have set the seed for random number generation, to ensure the reproducibility

```
Shape of the training dataset train_X: (199364, 29)
Shape of the testing dataset test_X: (85443, 29)
```

FindDefault-Credit card fraud detection

Applying Machine Learning Algorithm to Credit Card Dataset

- We will explore various machine learning algorithms to determine the most effective model for our binary classification problem.
- The task involves predicting one of the two class labels. We plan to assess the performance of different algorithms, such as Random Forest and Decision Tree to identify the most suitable solution for our specific problem.
- Our approach involves constructing Random Forest and Decision Tree classifiers to identify the most effective model.

Decision Tree Algorithm

The Decision Tree Algorithm is a supervised machine learning technique employed for both classification and regression tasks. Its objective is to create a training model capable of predicting the value of a target class variable. This is achieved by learning straightforward if-then-else decision rules derived from the patterns present in the training data.

Random Forest Algorithm

Random Forest is a supervised Machine Learning algorithm. It creates a "forest" out of an ensemble of "decision trees", which are normally trained using the "bagging" technique. The bagging method's basic principle is that combining different learning models improved the outcome. To get a more precise and reliable forecast, random forest creates several decision trees and merges them.

```
# Decision Tree
decision_tree = DecisionTreeClassifier()

# Random Forest
random_forest = RandomForestClassifier(n_estimators=100)
```

- Here we are creating a **RandomForestClassifier** with 100 trees in the forest.
- The large number of trees will generally lead to better performance, but it may also increase the training time.

Now we will check the score of the Decision Tree model

Decision Tree Score is: **99.91807403766254**

Random Forest Score is: **99.96254813150287**

```
# Printing the scores of the both classifiers
print("Decision Tree: ", round((decision_tree_score),4))
print("Random Forest: ", round((random_forest_score),4))

Decision Tree: 99.9181
Random Forest: 99.9625
```

FindDefault-Credit card fraud detection

The Random Forest classifier has a slight edge over the Decision Tree Classifier.


- We will create a function to print the metrics:

1. Accuracy_score
2. Precision_score
3. Confusion_matrix
4. Recall_score
5. F-1 score

```
Evaluation of Decision Tree Model:  
Accuracy: 0.9992  
Precision: 0.7062  
recall_score: 0.8309  
F1-Score: 0.7635
```

```
Evaluation of Random Forest Model:  
Accuracy: 0.9996  
Precision: 0.9561  
recall_score: 0.8015  
F1-Score: 0.872
```

Confusion Matrix



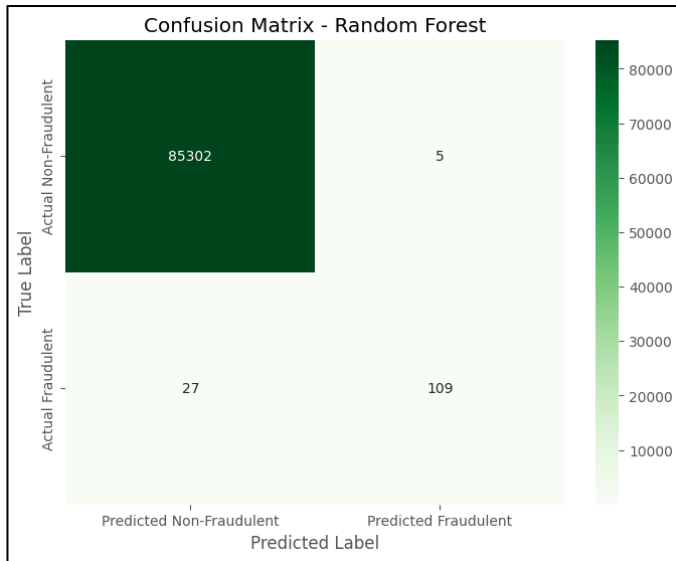
In-short summary:

1. The model correctly identified 107 fraudulent transactions.
2. It incorrectly identified 29 transactions as non-fraudulent.
3. It correctly identified 85265 non-fraudulent transactions.
4. It incorrectly identified 42 non-fraudulent transactions as fraudulent

We understand from the confusion matrix:

- Non-Fraudulent transactions:
 1. Correctly predicted as non-fraudulent(True Negative) 85265 transactions.
 2. Incorrectly predicted as fraudulent(False Positive) 42 transactions.
- Fraudulent Transactions:
 1. Incorrectly predicted as non-fraudulent(False Negative) : 29 transactions
 2. Correctly predicted as fraudulent(True Positive) : 107 transactions

FindDefault-Credit card fraud detection



In-short summary:

1. The model correctly identified 109 fraudulent transactions.
2. It incorrectly identified 27 transactions as non-fraudulent.
3. It correctly identified 85302 non-fraudulent transactions.
4. It incorrectly identified only 5 non-fraudulent transactions as fraudulent.

We understand from the confusion matrix:

- Non-Fraudulent transactions:
 1. Correctly predicted as non-fraudulent(True Negative) 85302 transactions.
 2. Incorrectly predicted as fraudulent(False Positive) 5 transactions.
- Fraudulent Transactions:
 1. Incorrectly predicted as non-fraudulent(False Negative) : 27 transactions
 2. Correctly predicted as fraudulent(True Positive) : 109 transactions

Class-Imbalance

- The Random Forest model works better than Decision Trees. In the presenece of a class-Imbalance issue, where genuine transactions account for over 99% of the dataset and credit card fraud transactions constitute only 0.17%.
- Training the model without addressing the imbalance can lead to biased predictions.
- Despite the apparent accuracy, such a model may not effectively capture the nuances of the minority class (fraud transactions) and may not generalize well to real-world situations.
- The class imbalance problem can be solved by various techniques. Oversampling is one of them.

FindDefault-Credit card fraud detection

- We will use the SMOT (Synthetic Minority Oversampling Technique, or SMOTE).It is the method of data augmentation for the minority class.

```
from imblearn.over_sampling import SMOTE

X_resampled, Y_resampled = SMOTE().fit_resample(X,Y)

print("Resampled shape of X: ",X_resampled.shape)
print("Resampled shape of Y: ",Y_resampled.shape)

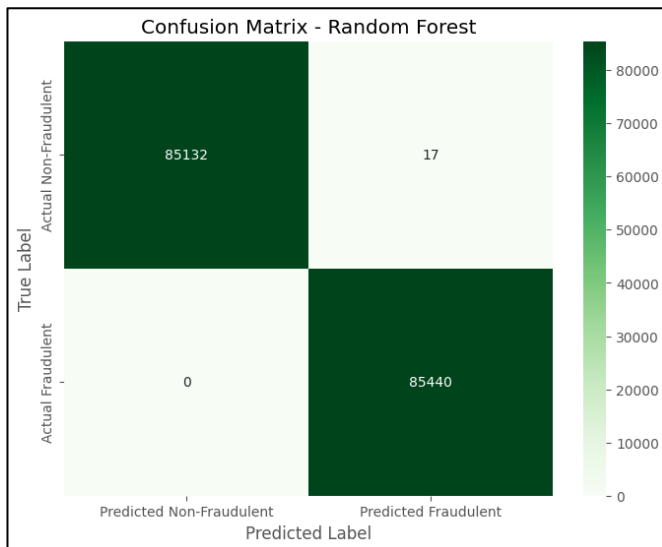
Resampled shape of X: (568630, 29)
Resampled shape of Y: (568630,)
```

```
predictions_resampled = rf_resampled.predict(test_X)
random_forest_score_resampled = rf_resampled.score(test_X, test_Y) * 100

print(random_forest_score_resampled)

99.99003452743142
```

Confusion matrix for resampled data



We understand from the confusion Matrix is:

- The model correctly identified 85440 fraudulent transactions.
- It incorrectly identified 0 transactions as non-fraudulent.
- It correctly identified 85131 non-fraudulent transactions.
- It incorrectly identified only 17 non-fraudulent transactions as fraudulent.

We understand from the confusion matrix:

- Non-Fraudulent transactions:
- Correctly predicted as non-fraudulent(True Negative) 85131 transactions.
- Incorrectly predicted as fraudulent(False Positive) 17 transactions.
- Fraudulent Transactions:
- Incorrectly predicted as non-fraudulent(False Negative) : 0 transactions

FindDefault-Credit card fraud detection

- Correctly predicted as fraudulent(True Positive) : 85440 transactions

```
Evaluation of Resampled Random Forest Model:  
Accuracy: 0.9999  
Precision: 0.9998  
recall_score: 1.0  
F1-Score: 0.9999
```

- We can see that our model performed much better than the previous Random Forest classifier without oversampling.
- We have applied the techniques to address the class imbalance issues and achieved an accuracy of more than 99%.

- We will import the pickle to dump the dataframe and the model for the model deployment as the future scope.

```
import pickle  
pickle.dump(df,open('df.pkl','wb'))  
pickle.dump(rf_resampled,open('rf_resampled.pkl','wb'))
```